

Position / Direction / Object Dim Analysis — Vision Encoder output (SigLIP, mean-pooled)
Top-50 dims per task × attribute → cross-task overlap
R2R_4way_1500 (6000 samples/task, 3 models → Vision Encoder frozen same across)

